

Chapter 10: Deep Learning: Training & Optimisation

Learning Objectives

By the end of this chapter, you will be able to:

- **Master** deep neural network architectures and training techniques
- **Understand** optimization algorithms and their underlying mathematical principles
- **Implement** key deep learning components from scratch and with frameworks
- **Apply** best practices for model development and troubleshooting
- **Compare** biological and artificial optimization approaches

10.1 Neural Network Fundamentals

Deep learning has revolutionized AI by enabling models to learn hierarchical representations directly from data. At its core, deep learning is built on neural networks with multiple layers that progressively extract higher-level features.

$$h = \sigma(Wx + b)$$

Neural Network Architecture

Multilayer Perceptron with Hidden Layers

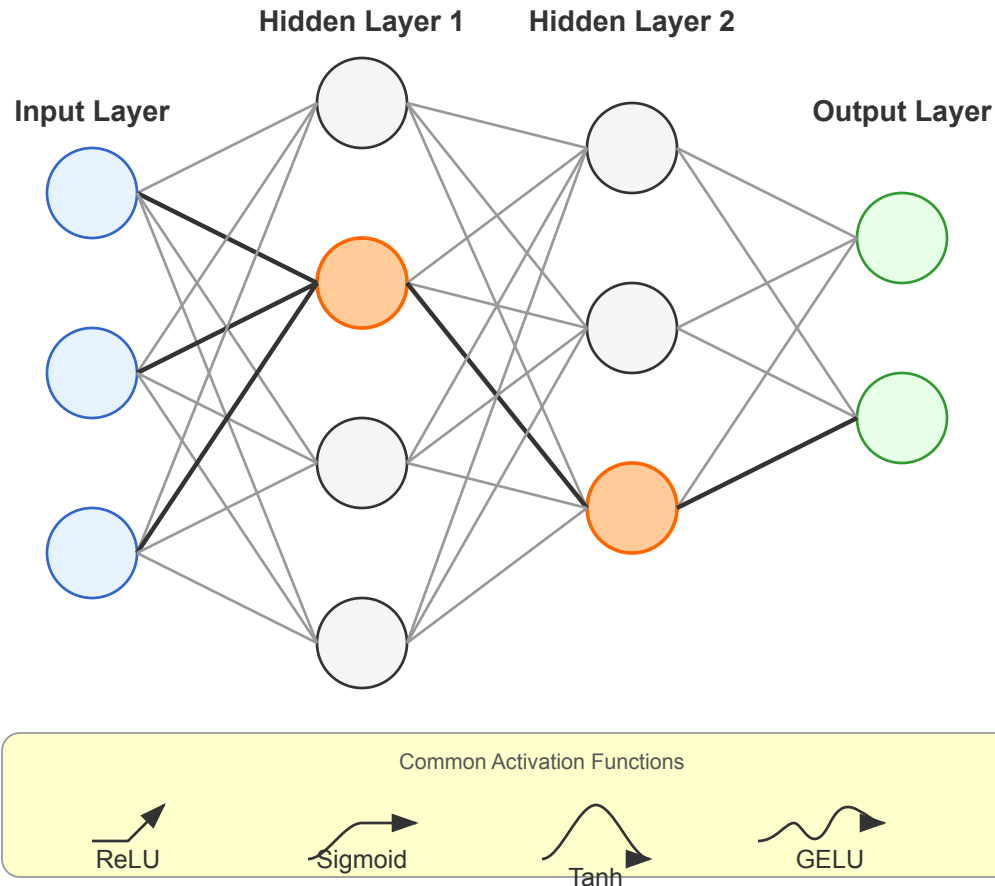


Figure 10.1: A multilayer perceptron with two hidden layers, showing how neurons connect between layers and activate through non-linear functions.

10.1.1 Multilayer Perceptrons

Multilayer perceptrons (MLPs) are the foundational architecture of deep learning, consisting of an input layer, one or more hidden layers, and an output layer.

```

import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim

# Define a simple MLP using PyTorch
class SimpleMLP(nn.Module):
    def __init__(self, input_size, hidden_sizes, output_size):
        """
        A basic multilayer perceptron implementation.

        Args:
            input_size: Number of input features
            hidden_sizes: List of hidden layer sizes
            output_size: Number of output units
        """
        super(SimpleMLP, self).__init__()

        # Create a list to hold all layers
        layers = []

        # Input layer to first hidden layer
        layers.append(nn.Linear(input_size, hidden_sizes[0]))
        layers.append(nn.ReLU())

        # Hidden layers
        for i in range(len(hidden_sizes) - 1):
            layers.append(nn.Linear(hidden_sizes[i], hidden_sizes[i+1]))
            layers.append(nn.ReLU())

        # Final layer
        layers.append(nn.Linear(hidden_sizes[-1], output_size))

        # Combine all layers into a sequential model
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        """Forward pass through the network."""
        return self.model(x)

# Create an example MLP
input_size = 10
hidden_sizes = [128, 64, 32]
output_size = 2
model = SimpleMLP(input_size, hidden_sizes, output_size)

# Display model architecture
print(model)

```

The power of MLPs comes from their ability to approximate any continuous function with sufficient neurons in the hidden layers (universal approximation theorem).

10.1.2 Activation Functions

Activation functions introduce non-linearity into neural networks, allowing them to learn complex patterns.

```
def plot_activation_functions():
    """Plot common activation functions used in deep learning."""
    # Generate input values
    x = np.linspace(-5, 5, 1000)

    # Calculate activation function outputs
    sigmoid = 1 / (1 + np.exp(-x))
    tanh = np.tanh(x)
    relu = np.maximum(0, x)
    leaky_relu = np.where(x > 0, x, 0.1 * x)
    elu = np.where(x > 0, x, np.exp(x) - 1)

    # Create plot
    plt.figure(figsize=(12, 8))

    plt.plot(x, sigmoid, label='Sigmoid')
    plt.plot(x, tanh, label='Tanh')
    plt.plot(x, relu, label='ReLU')
    plt.plot(x, leaky_relu, label='Leaky ReLU')
    plt.plot(x, elu, label='ELU')

    plt.grid(True)
    plt.legend()
    plt.title('Common Activation Functions')
    plt.xlabel('Input')
    plt.ylabel('Output')
    plt.axhline(y=0, color='k', linestyle='--', alpha=0.3)
    plt.axvline(x=0, color='k', linestyle='--', alpha=0.3)

    plt.tight_layout()
    return plt
```

Common activation functions include:

- **ReLU (Rectified Linear Unit):** $f(x) = \max(0, x)$
 - Pros: Fast computation, reduces vanishing gradient problem
 - Cons: "Dying ReLU" problem (neurons can get stuck)

- **Leaky ReLU:** $f(x) = \max(\alpha x, x)$ where α is a small constant
 - Pros: Addresses dying ReLU problem
 - Cons: Performance improvement is often marginal
- **Sigmoid:** $f(x) = \frac{1}{1+e^{-x}}$
 - Pros: Outputs between 0 and 1, useful for binary classification
 - Cons: Vanishing gradient problem for extreme inputs
- **Tanh (Hyperbolic Tangent):** $f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
 - Pros: Zero-centered, useful in recurrent networks
 - Cons: Still suffers from vanishing gradient
- **GELU (Gaussian Error Linear Unit):** $f(x) = x \cdot \Phi(x)$ where Φ is the CDF of the standard normal distribution
 - Pros: Smooth, better performance in transformers
 - Cons: More computationally expensive

Recent architectures like transformers commonly use GELU, while CNNs often use ReLU or its variants.

10.1.3 Backpropagation Algorithm

Backpropagation is the cornerstone algorithm for training neural networks, efficiently computing gradients through the chain rule.

```

def manual_backpropagation_example():
    """
    Demonstrate backpropagation with a simple 2-layer network.
    """
    # Simple network: Input -> Hidden (2 neurons) -> Output
    # Forward pass
    def sigmoid(x):
        return 1 / (1 + np.exp(-x))

    def sigmoid_derivative(x):
        return x * (1 - x)

    # Network parameters
    input_size = 3
    hidden_size = 2
    output_size = 1

    # Inputs and target
    X = np.array([[0.1, 0.2, 0.3]])
    y_true = np.array([[0.7]])

    # Initialize weights and biases
    np.random.seed(42)
    W1 = np.random.randn(input_size, hidden_size)
    b1 = np.zeros((1, hidden_size))
    W2 = np.random.randn(hidden_size, output_size)
    b2 = np.zeros((1, output_size))

    # Forward pass
    hidden_input = np.dot(X, W1) + b1
    hidden_output = sigmoid(hidden_input)
    final_input = np.dot(hidden_output, W2) + b2
    y_pred = sigmoid(final_input)

    # Calculate loss
    loss = 0.5 * np.sum((y_pred - y_true) ** 2)

    # Backpropagation
    # Output layer error
    output_error = y_pred - y_true
    output_delta = output_error * sigmoid_derivative(y_pred)

    # Hidden layer error
    hidden_error = np.dot(output_delta, W2.T)
    hidden_delta = hidden_error * sigmoid_derivative(hidden_output)

    # Update weights and biases
    learning_rate = 0.1
    W2 -= learning_rate * np.dot(hidden_output.T, output_delta)
    b2 -= learning_rate * np.sum(output_delta, axis=0, keepdims=True)
    W1 -= learning_rate * np.dot(X.T, hidden_delta)
    b1 -= learning_rate * np.sum(hidden_delta, axis=0, keepdims=True)

```

```
return {  
    'Initial prediction': y_pred[0][0],  
    'Target': y_true[0][0],  
    'Loss': loss,  
    'Output delta': output_delta[0][0],  
    'Hidden delta': hidden_delta[0]  
}
```

The backpropagation algorithm:

1. **Forward Pass:** Compute outputs of all neurons from input to output
2. **Error Calculation:** Compare network output with target to compute error
3. **Backward Pass:** Propagate error backward to assign “responsibility” to each parameter
4. **Parameter Update:** Adjust weights and biases using calculated gradients

While modern deep learning frameworks handle these calculations automatically through automatic differentiation, understanding backpropagation is crucial for developing intuition about neural network training.

10.1.4 Vanishing/Exploding Gradients

As neural networks get deeper, the problem of vanishing or exploding gradients becomes more severe:

```
def demonstrate_gradient_problems():
    """
    Visualize vanishing/exploding gradient problems in deep networks.
    """
    depths = list(range(1, 21))

    # Vanishing gradient with sigmoid
    vanishing_grads = [0.25 ** d for d in depths]

    # Exploding gradient with poor initialization
    exploding_grads = [1.5 ** d for d in depths]

    # Plot
    plt.figure(figsize=(10, 6))
    plt.semilogy(depths, vanishing_grads, 'b-', label='Vanishing Gradient (sigmoid)')
    plt.semilogy(depths, exploding_grads, 'r-', label='Exploding Gradient (poor initialization)')
    plt.semilogy(depths, [0.1] * len(depths), 'g--', label='Stable Gradient (with proper init)')

    plt.xlabel('Network Depth (layers)')
    plt.ylabel('Gradient Magnitude (log scale)')
    plt.title('Vanishing and Exploding Gradients in Deep Networks')
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    return plt
```

Solutions to these gradient problems include:

- Careful weight initialization (e.g., He, Xavier/Glorot)
- Batch normalization
- Residual connections
- Gradient clipping
- Using activation functions that don't saturate (e.g., ReLU)

10.2 Optimization Techniques

Training deep neural networks requires effective optimization algorithms. The choice of optimizer significantly impacts training speed and model performance.

Optimization Algorithms Comparison

Convergence paths in loss landscape




Figure 10.2: Comparison of different optimization algorithms showing convergence paths in a loss landscape. Modern approaches like Adam and RMSprop often converge faster and more reliably than vanilla SGD.

10.2.1 Stochastic Gradient Descent

Stochastic Gradient Descent (SGD) is the most fundamental optimization algorithm for neural networks:

```

def sgd_optimizer_example():
    """Implement basic SGD and variants."""
    # Generate dummy data
    np.random.seed(42)
    X = np.random.randn(1000, 10)
    w_true = np.random.randn(10, 1)
    y = X @ w_true + 0.1 * np.random.randn(1000, 1)

    # Initialize parameters
    w = np.zeros((10, 1))

    # SGD parameters
    learning_rate = 0.01
    batch_size = 32
    epochs = 100

    # Training loop
    losses = []
    for epoch in range(epochs):
        # Shuffle data
        indices = np.random.permutation(len(X))
        X_shuffled = X[indices]
        y_shuffled = y[indices]

        epoch_losses = []
        # Process mini-batches
        for i in range(0, len(X), batch_size):
            X_batch = X_shuffled[i:i+batch_size]
            y_batch = y_shuffled[i:i+batch_size]

            # Forward pass
            y_pred = X_batch @ w
            loss = np.mean((y_pred - y_batch) ** 2)
            epoch_losses.append(loss)

            # Backward pass (compute gradient)
            grad = 2 * X_batch.T @ (y_pred - y_batch) / batch_size

            # Update parameters
            w = w - learning_rate * grad

        losses.append(np.mean(epoch_losses))

    return {
        'weights': w,
        'true_weights': w_true,
        'final_loss': losses[-1],
        'loss_history': losses
    }

```

SGD variants include:

- **Batch Gradient Descent:** Uses the entire dataset per update
- **Mini-batch SGD:** Uses small batches (typically 32-256 samples)
- **Online SGD:** Updates using one sample at a time

10.2.2 Momentum and Adaptive Methods

Modern optimizers build upon SGD by introducing momentum or adaptive learning rates:

```

def compare_optimizers():
    """Compare convergence speed of different optimizers."""
    import torch
    import torch.nn as nn
    import torch.optim as optim

    # Define a simple problem
    X = torch.randn(1000, 20)
    y = torch.randn(1000, 1)

    # Define model
    model = nn.Sequential(
        nn.Linear(20, 64),
        nn.ReLU(),
        nn.Linear(64, 1)
    )

    # Loss function
    criterion = nn.MSELoss()

    # Define optimizers
    optimizers = {
        'SGD': optim.SGD(model.parameters(), lr=0.01),
        'SGD+Momentum': optim.SGD(model.parameters(), lr=0.01, momentum=0.9),
        'Adam': optim.Adam(model.parameters(), lr=0.01),
        'RMSprop': optim.RMSprop(model.parameters(), lr=0.01),
        'AdamW': optim.AdamW(model.parameters(), lr=0.01, weight_decay=1e-4)
    }

    # Training loops
    results = {}

    for name, optimizer in optimizers.items():
        # Reset model
        model = nn.Sequential(
            nn.Linear(20, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )

        # Train
        losses = []
        for epoch in range(100):
            optimizer.zero_grad()
            outputs = model(X)
            loss = criterion(outputs, y)
            losses.append(loss.item())
            loss.backward()
            optimizer.step()

        results[name] = losses

    # Plot

```

```

plt.figure(figsize=(10, 6))
for name, loss_history in results.items():
    plt.plot(loss_history, label=name)

plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Optimizer Convergence Comparison')
plt.legend()
plt.yscale('log')
plt.grid(True)

plt.tight_layout()
return plt

```

Key optimization algorithms include:

- **SGD with Momentum:** Adds a fraction of the previous update to the current one, helping to escape local minima and accelerate convergence
 - $v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$
 - $\theta = \theta - v_t$
- **Nesterov Accelerated Gradient:** Computes gradient at the “looked-ahead” position for better convergence
 - $v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$
 - $\theta = \theta - v_t$
- **AdaGrad:** Adapts learning rates per-parameter based on historical gradients
 - $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t$
- **RMSprop:** Modifies AdaGrad to better handle non-convex functions by using an exponentially weighted moving average
 - $E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta) g_t^2$
 - $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$
- **Adam:** Combines momentum and RMSprop ideas for robust performance
 - $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$ (momentum)
 - $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$ (RMSprop)
 - $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$ (bias correction)
 - $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$

Adam is currently the most widely used optimizer due to its robustness across different architectures and datasets.

10.2.3 Learning Rate Schedules

Learning rate scheduling can significantly improve training outcomes:

```

def learning_rate_schedules():
    """Visualize common learning rate schedules."""
    epochs = np.arange(1, 101)

    # Constant
    constant_lr = [0.1] * 100

    # Step decay
    step_lr = [0.1 * (0.1 ** (e // 30)) for e in epochs]

    # Exponential decay
    exp_lr = [0.1 * np.exp(-0.03 * e) for e in epochs]

    # Cosine annealing
    cosine_lr = [0.1 * (1 + np.cos(np.pi * e / 100)) / 2 for e in epochs]

    # Linear warmup + cosine decay
    warmup = 10
    warmup_cosine_lr = []
    for e in epochs:
        if e <= warmup:
            lr = 0.1 * e / warmup
        else:
            lr = 0.1 * (1 + np.cos(np.pi * (e - warmup) / (100 - warmup))) / 2
        warmup_cosine_lr.append(lr)

    # Plot
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, constant_lr, label='Constant')
    plt.plot(epochs, step_lr, label='Step Decay')
    plt.plot(epochs, exp_lr, label='Exponential Decay')
    plt.plot(epochs, cosine_lr, label='Cosine Annealing')
    plt.plot(epochs, warmup_cosine_lr, label='Warmup + Cosine')

    plt.xlabel('Epoch')
    plt.ylabel('Learning Rate')
    plt.title('Learning Rate Schedules')
    plt.legend()
    plt.yscale('log')
    plt.grid(True)

    plt.tight_layout()
    return plt

```

Popular learning rate schedules include:

- **Step Decay:** Reduces learning rate by a factor after a set number of epochs
- **Exponential Decay:** Continuously decreases learning rate using an exponential function
- **Cosine Annealing:** Smoothly decreases learning rate following a cosine curve

- **Cyclic Learning Rates:** Cycles between lower and upper learning rate bounds
- **One-Cycle Policy:** Increases learning rate to a maximum, then decreases it
- **Warmup + Decay:** Gradually increases learning rate during initial epochs, then decays

10.2.4 Second-order Methods

While first-order methods like SGD use only gradient information, second-order methods incorporate curvature information:


```

def second_order_methods():
    """Compare first and second-order optimization methods."""
    # Generate a 2D quadratic function with conditioning issues
    def f(x, y):
        return 0.01 * x**2 + 5 * y**2

    def grad_f(x, y):
        return np.array([0.02 * x, 10 * y])

    def hessian_f(x, y):
        return np.array([[0.02, 0], [0, 10]])

    # Starting point
    x0, y0 = 10.0, 2.0

    # SGD trajectory
    sgd_path = [(x0, y0)]
    x, y = x0, y0
    lr = 0.1
    for _ in range(20):
        g = grad_f(x, y)
        x -= lr * g[0]
        y -= lr * g[1]
        sgd_path.append((x, y))

    # Newton's method trajectory
    newton_path = [(x0, y0)]
    x, y = x0, y0
    for _ in range(5): # Usually converges in fewer steps
        g = grad_f(x, y)
        H = hessian_f(x, y)
        H_inv = np.linalg.inv(H)
        update = H_inv @ g
        x -= update[0]
        y -= update[1]
        newton_path.append((x, y))

    # Plot
    x_range = np.linspace(-10, 10, 100)
    y_range = np.linspace(-2, 2, 100)
    X, Y = np.meshgrid(x_range, y_range)
    Z = f(X, Y)

    plt.figure(figsize=(10, 8))

    # Contour plot
    plt.contour(X, Y, Z, 20, cmap='viridis', alpha=0.6)

    # Plot paths
    sgd_path = np.array(sgd_path)
    newton_path = np.array(newton_path)

    plt.plot(sgd_path[:, 0], sgd_path[:, 1], 'r.-', label='SGD', linewidth=2, mar

```

```
plt.plot(newton_path[:, 0], newton_path[:, 1], 'b.-', label="Newton's Method")

plt.xlabel('x')
plt.ylabel('y')
plt.title("Comparison of First-Order vs. Second-Order Methods")
plt.legend()
plt.grid(True)

plt.tight_layout()
return plt
```

Second-order methods include:

- **Newton's Method:** Uses the Hessian matrix (second derivatives) for updates
 - $\theta_{t+1} = \theta_t - H^{-1}(\theta_t) \nabla_{\theta} J(\theta_t)$
 - Pros: Fast convergence near optimum
 - Cons: Expensive Hessian computation and inversion
- **Quasi-Newton Methods** (e.g., BFGS, L-BFGS): Approximate the Hessian
 - Pros: Faster than Newton's method
 - Cons: Still too expensive for large neural networks
- **Natural Gradient Descent:** Uses the Fisher information matrix
 - Pros: Invariant to reparameterization
 - Cons: Computation and storage requirements

While second-order methods offer theoretical advantages, their computational requirements generally make them impractical for deep learning. However, approximations like K-FAC (Kronecker-Factored Approximate Curvature) are being explored to make second-order information more accessible.

10.3 Regularization Strategies

Regularization helps prevent overfitting by constraining the model's capacity or adding noise to the training process.

Regularization Techniques in Deep Learning

Methods to prevent overfitting and improve generalization

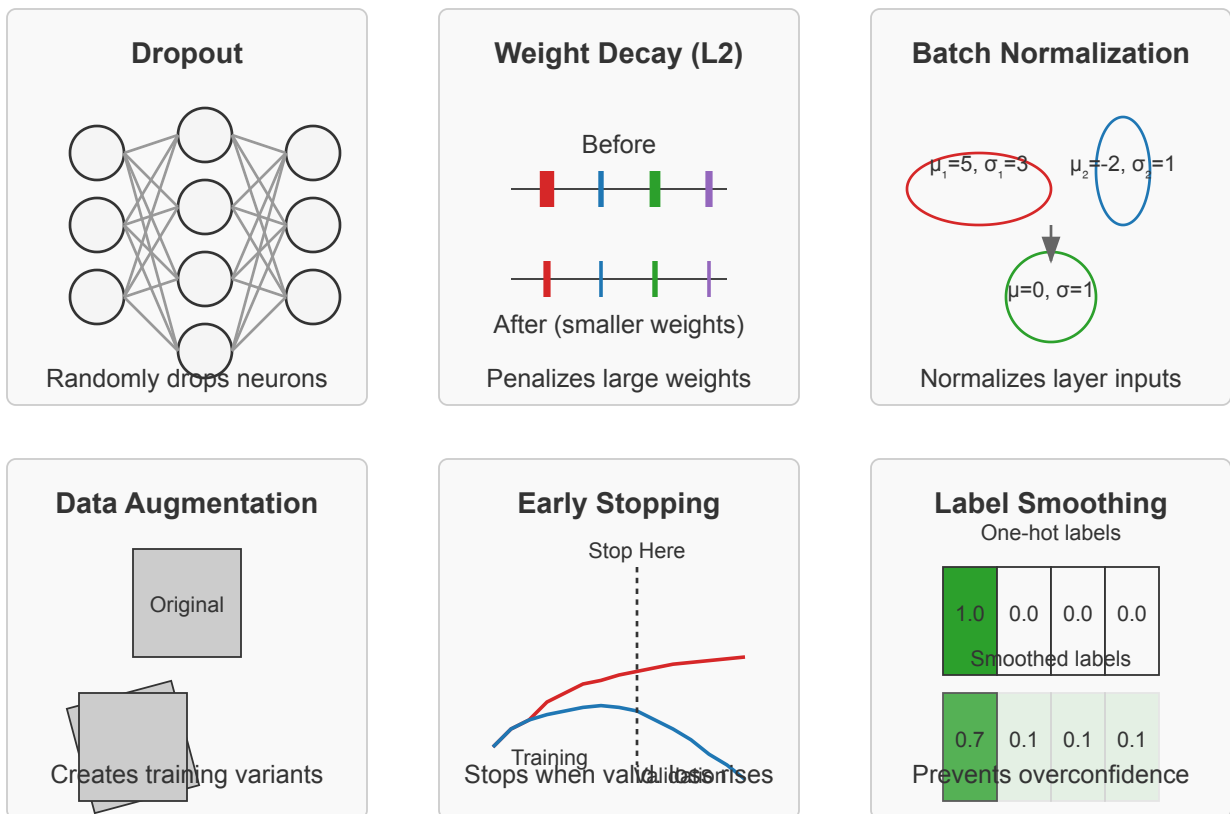


Figure 10.3: Common regularization methods in deep learning, including dropout, weight decay (L2), batch normalization, data augmentation, early stopping, and label smoothing.

10.3.1 Dropout and Batch Normalization

```
def dropout_example():
    """Implement and visualize dropout."""
    import torch
    import torch.nn as nn

    # Define a model with dropout
    class MLPWithDropout(nn.Module):
        def __init__(self, dropout_rate=0.5):
            super().__init__()
            self.fc1 = nn.Linear(784, 256)
            self.dropout1 = nn.Dropout(dropout_rate)
            self.fc2 = nn.Linear(256, 128)
            self.dropout2 = nn.Dropout(dropout_rate)
            self.fc3 = nn.Linear(128, 10)

        def forward(self, x):
            x = torch.relu(self.fc1(x))
            x = self.dropout1(x)
            x = torch.relu(self.fc2(x))
            x = self.dropout2(x)
            x = self.fc3(x)
            return x

    # Create a toy example for visualization
    model = MLPWithDropout(dropout_rate=0.5)

    # Generate random activations for demonstration
    activations = torch.rand(1, 256)

    # Apply dropout with different rates
    dropout_rates = [0.0, 0.3, 0.5, 0.7]
    results = []

    for rate in dropout_rates:
        dropout = nn.Dropout(rate)
        # Set model to training mode
        dropout.train()
        # Apply dropout
        dropped_activations = dropout(activations)
        results.append((rate, dropped_activations[0].detach().numpy()))

    # Visualize
    fig, axes = plt.subplots(len(dropout_rates), 1, figsize=(10, 8))

    for i, (rate, acts) in enumerate(results):
        axes[i].bar(range(50), acts[:50], alpha=0.7) # Show first 50 units
        axes[i].set_title(f'Dropout Rate: {rate}')
        axes[i].set_ylim(0, 1.5) # Account for scaling during training
        axes[i].grid(True, alpha=0.3)

    plt.tight_layout()
```

```

return fig

def batch_norm_example():
    """Demonstrate batch normalization effect."""
    import torch
    import torch.nn as nn

    # Create random activations
    np.random.seed(42)
    # Poorly scaled/shifted activations
    activations = np.random.randn(100, 32) * 10 + 5

    # Apply batch normalization
    bn = nn.BatchNorm1d(32)
    normalized = bn(torch.tensor(activations, dtype=torch.float32))
    normalized = normalized.detach().numpy()

    # Visualize
    fig, axes = plt.subplots(2, 1, figsize=(10, 8))

    # Plot raw activation distribution
    for i in range(5): # Show first 5 features
        axes[0].hist(activations[:, i], alpha=0.3, bins=20, label=f'Feature {i+1}')
    axes[0].set_title('Before Batch Normalization')
    axes[0].grid(True, alpha=0.3)
    axes[0].legend()

    # Plot normalized activation distribution
    for i in range(5):
        axes[1].hist(normalized[:, i], alpha=0.3, bins=20, label=f'Feature {i+1}')
    axes[1].set_title('After Batch Normalization')
    axes[1].grid(True, alpha=0.3)
    axes[1].legend()

    plt.tight_layout()
    return fig

```

Dropout stochastically zeroes activations during training, forcing the network to learn redundant representations:

1. During training: Each neuron is kept with probability p (typically 0.5 to 0.8)
2. During inference: All neurons are used, but outputs are scaled by p

Batch Normalization normalizes activations within a mini-batch, making training more stable:

1. Normalize: $\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$
2. Scale and shift: $y_i = \gamma \hat{x}_i + \beta$

Benefits include:

- Reduced internal covariate shift
- Improved gradient flow
- Regularization effect
- Reduced sensitivity to initialization

10.3.2 Weight Decay and Early Stopping

```
def weight_decay_visualization():
    """Visualize the effect of weight decay on model complexity."""
    from sklearn.linear_model import Ridge

    # Generate synthetic data
    np.random.seed(42)
    X = np.sort(np.random.rand(100, 1) * 6 - 3, axis=0)
    y = np.sin(X.ravel()) + np.random.normal(0, 0.1, X.shape[0])

    # Fit with different regularization strengths
    alphas = [0, 0.001, 0.01, 0.1, 1.0]
    degrees = 10 # polynomial degree

    X_plot = np.linspace(-3, 3, 1000).reshape(-1, 1)

    plt.figure(figsize=(12, 8))

    for i, alpha in enumerate(alphas):
        # Create polynomial features
        from sklearn.preprocessing import PolynomialFeatures
        from sklearn.pipeline import make_pipeline

        model = make_pipeline(
            PolynomialFeatures(degrees),
            Ridge(alpha=alpha)
        )

        model.fit(X, y)
        y_plot = model.predict(X_plot)

        plt.subplot(len(alphas), 1, i+1)
        plt.scatter(X, y, color='navy', s=30, marker='o', label="Training data")
        plt.plot(X_plot, y_plot, color='red', label="Model")
        plt.plot(X_plot, np.sin(X_plot.ravel()), color='green', label="True funct")
        plt.title(f"Weight Decay (L2):  $\alpha = \{alpha\}$ ")
        plt.ylim((-1.5, 1.5))
        plt.legend()

    plt.tight_layout()
    return plt

def plot_early_stopping():
    """Visualize early stopping based on validation performance."""
    # Simulate training and validation losses
    epochs = np.arange(1, 101)

    # Training loss (continues to decrease)
    train_loss = 1.0 / (0.1 * epochs + 1.0) + 0.1

    # Validation loss (starts increasing after a while)
    val_loss = 1.0 / (0.1 * epochs + 1.0) + 0.1 + 0.05 * np.maximum(0, epochs - 4
```

```

# Add noise
np.random.seed(42)
train_loss += np.random.normal(0, 0.02, len(epochs))
val_loss += np.random.normal(0, 0.03, len(epochs))

# Determine early stopping point
patience = 10
best_val_loss = float('inf')
best_epoch = 0
stop_epoch = 0

for i, loss in enumerate(val_loss):
    if loss < best_val_loss:
        best_val_loss = loss
        best_epoch = i
    elif i > best_epoch + patience:
        stop_epoch = i
        break

# Plot
plt.figure(figsize=(10, 6))
plt.plot(epochs, train_loss, 'b-', label='Training Loss')
plt.plot(epochs, val_loss, 'r-', label='Validation Loss')

# Mark early stopping point
if stop_epoch > 0:
    plt.axvline(x=stop_epoch, color='g', linestyle='--', label=f'Early Stoppi

# Mark best validation point
plt.axvline(x=best_epoch, color='m', linestyle=':', label=f'Best Validation (

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Early Stopping Based on Validation Loss')
plt.legend()
plt.grid(True)

plt.tight_layout()
return plt

```

Weight Decay (L2 regularization) adds a penalty term to the loss function proportional to the squared weights:

$$L_{reg} = L_{original} + \lambda \sum_i w_i^2$$

This encourages the model to use smaller weights, reducing model complexity and preventing overfitting.

Early Stopping halts training when performance on a validation set stops improving:

1. Monitor validation performance at regular intervals
2. Save the model when it achieves the best validation performance
3. Stop training after a predefined number of epochs without improvement (patience)
4. Restore the best model from the saved checkpoint

Early stopping effectively limits the model's capacity by restricting the number of optimization steps.

10.3.3 Data Augmentation

```
def data_augmentation_example():
    """Demonstrate common data augmentation techniques."""
    try:
        from PIL import Image
        import torchvision.transforms as transforms
        import torchvision.transforms.functional as TF

        # Create a sample image (a simple placeholder)
        img = Image.new('RGB', (300, 200), color=(73, 109, 137))

        # Define augmentations
        augmentations = [
            ('Original', lambda x: x),
            ('Horizontal Flip', TF.hflip),
            ('Rotation (30°)', lambda x: TF.rotate(x, 30)),
            ('Random Crop', lambda x: TF.crop(x, 50, 50, 150, 100)),
            ('Color Jitter', lambda x: transforms.ColorJitter(brightness=0.5, con
            ('Random Erasing', lambda x: transforms.RandomErasing(p=1.0, scale=(0
        ]

        # Apply and visualize
        fig, axes = plt.subplots(2, 3, figsize=(12, 8))
        axes = axes.flatten()

        for i, (name, aug_fn) in enumerate(augmentations):
            if name == 'Random Erasing':
                # Special case for random erasing which expects a tensor
                axes[i].imshow(aug_fn.permute(1, 2, 0))
            else:
                axes[i].imshow(aug_fn(img))
                axes[i].set_title(name)
                axes[i].axis('off')

        plt.tight_layout()
        return fig

    except ImportError:
        # If PIL or torchvision not available, return a text figure
        fig, ax = plt.subplots(figsize=(10, 6))
        ax.text(0.5, 0.5, "Data Augmentation Techniques:\n\n" +
            "• Horizontal/Vertical Flips\n" +
            "• Random Rotations\n" +
            "• Random Crops\n" +
            "• Color Jitter (brightness, contrast, saturation)\n" +
            "• Random Erasing\n" +
            "• Cutout/CutMix/MixUp\n" +
            "• Elastic Transformations",
            horizontalalignment='center',
            verticalalignment='center',
            fontsize=14)
```

```
ax.axis('off')  
return fig
```

Data augmentation artificially increases the size of the training set by applying transformations to the original data:

- **Image Augmentations:** Flips, rotations, crops, color adjustments, random erasing
- **Advanced Techniques:** Mixup (blend images and labels), CutMix (patch replacement), AugMix (augmentation chains)
- **Text Augmentations:** Synonym replacement, word insertion/deletion, back-translation
- **Audio Augmentations:** Time stretching, pitch shifting, noise addition, spectrogram masking

Benefits include:

- Improved generalization
- Robustness to variations
- Reduced overfitting
- Better class balance

10.3.4 Label Smoothing

```
def label_smoothing_example():
    """Demonstrate the effect of label smoothing on model confidence."""
    # Calculate softmax probabilities
    def softmax(x):
        e_x = np.exp(x - np.max(x))
        return e_x / e_x.sum()

    # Loss functions
    def cross_entropy(probs, label, epsilon=0.0):
        """Cross entropy with optional label smoothing."""
        n_classes = len(probs)

        # Create one-hot encoding
        targets = np.zeros_like(probs)
        targets[label] = 1.0

        if epsilon > 0:
            # Apply label smoothing
            targets = (1 - epsilon) * targets + epsilon / n_classes

        # Compute loss
        return -np.sum(targets * np.log(probs + 1e-9))

    # Generate some logits
    logits_correct = np.array([10.0, 2.0, 1.0, 0.5, 0.1]) # Strongly predicting
    logits_wrong = np.array([2.0, 10.0, 1.0, 0.5, 0.1]) # Strongly predicting

    probs_correct = softmax(logits_correct)
    probs_wrong = softmax(logits_wrong)

    # Class 0 is the true label
    true_label = 0

    # Compare losses with and without label smoothing
    smoothing_values = [0.0, 0.1, 0.2]
    results = []

    for epsilon in smoothing_values:
        loss_correct = cross_entropy(probs_correct, true_label, epsilon)
        loss_wrong = cross_entropy(probs_wrong, true_label, epsilon)

        results.append({
            'epsilon': epsilon,
            'loss_correct': loss_correct,
            'loss_wrong': loss_wrong,
            'ratio': loss_wrong / loss_correct
        })

    # Create comparison plot
    fig, axes = plt.subplots(1, 2, figsize=(12, 6))
```

```

# Plot probabilities
bar_positions = np.arange(5)
axes[0].bar(bar_positions - 0.2, probs_correct, width=0.4, label='Correct Pre
axes[0].bar(bar_positions + 0.2, probs_wrong, width=0.4, label='Incorrect Pre
axes[0].set_xticks(bar_positions)
axes[0].set_xticklabels([f'Class {i}' for i in range(5)])
axes[0].set_ylabel('Probability')
axes[0].set_title('Model Predictions')
axes[0].legend()

# Plot loss comparisons
eps_values = [r['epsilon'] for r in results]
correct_losses = [r['loss_correct'] for r in results]
wrong_losses = [r['loss_wrong'] for r in results]
ratios = [r['ratio'] for r in results]

ax1 = axes[1]
ax1.plot(eps_values, correct_losses, 'b-o', label='Loss (Correct)')
ax1.plot(eps_values, wrong_losses, 'r-o', label='Loss (Incorrect)')
ax1.set_xlabel('Label Smoothing ( $\epsilon$ )')
ax1.set_ylabel('Loss Value')
ax1.set_title('Effect of Label Smoothing on Loss')
ax1.legend(loc='upper left')

ax2 = ax1.twinx()
ax2.plot(eps_values, ratios, 'g--s', label='Loss Ratio (Wrong/Correct)')
ax2.set_ylabel('Loss Ratio', color='g')
ax2.tick_params(axis='y', labelcolor='g')
ax2.legend(loc='upper right')

plt.tight_layout()
return fig

```

Label smoothing replaces one-hot encoded targets with “soft” targets:

$$y_i = \begin{cases} 1 - \epsilon + \epsilon/K & \text{if } i = \text{true class} \\ \epsilon/K & \text{otherwise} \end{cases}$$

where ϵ is the smoothing parameter and K is the number of classes.

Benefits include:

- Prevents overconfidence
- Improves generalization
- Provides regularization
- Aligns better with inherent data ambiguity

10.4 Advanced Architectures

Modern deep learning has evolved sophisticated architectures for different domains.

Advanced Neural Network Architectures

Key design patterns in modern deep learning

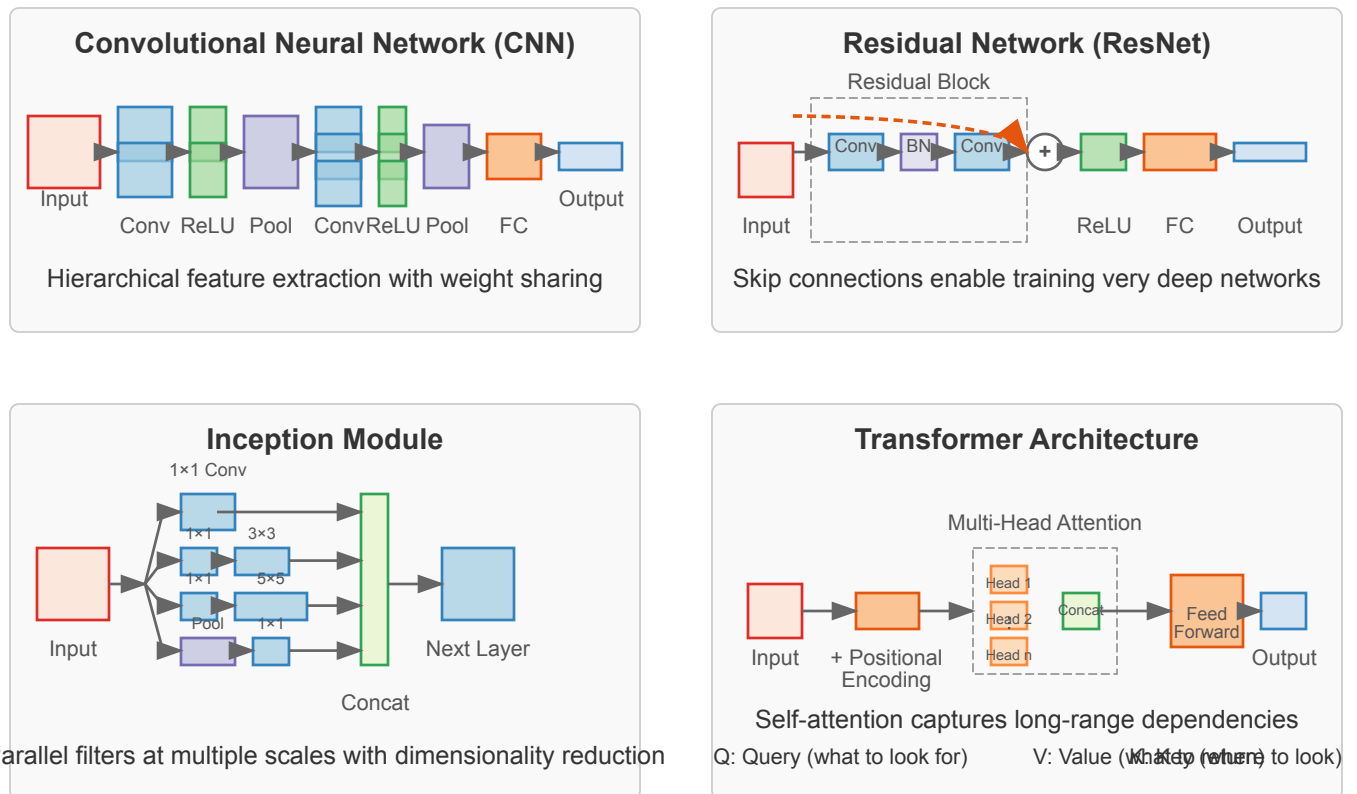


Figure 10.4: Key architectural patterns in modern deep learning, including CNNs, ResNets, Inception modules, and Transformer blocks, each addressing specific model design challenges.

10.4.1 Convolutional Neural Networks

```
def visualize_cnn_architecture():
    """Visualize a basic CNN architecture."""
    fig, ax = plt.subplots(figsize=(12, 5))

    # Define architecture components
    components = [
        {"name": "Input", "shape": (64, 64, 3), "x": 0.1, "width": 0.1},
        {"name": "Conv 3x3\n64 filters", "shape": (32, 32, 64), "x": 0.25, "width": 0.1},
        {"name": "Conv 3x3\n128 filters", "shape": (16, 16, 128), "x": 0.4, "width": 0.1},
        {"name": "MaxPool\n2x2", "shape": (8, 8, 128), "x": 0.55, "width": 0.07},
        {"name": "Flatten", "shape": "(8192,)", "x": 0.67, "width": 0.05},
        {"name": "Dense\n512 units", "shape": "(512,)", "x": 0.77, "width": 0.08},
        {"name": "Dense\n10 units", "shape": "(10,)", "x": 0.9, "width": 0.05}
    ]

    # Draw boxes
    for i, comp in enumerate(components):
        color = plt.cm.viridis(i / len(components))
        height = min(0.2 + 0.05 * i, 0.5)

        # Draw component box
        rect = plt.Rectangle(
            (comp["x"], 0.5 - height/2),
            comp["width"], height,
            facecolor=color, alpha=0.7, edgecolor='black'
        )
        ax.add_patch(rect)

        # Add labels
        ax.text(comp["x"] + comp["width"]/2, 0.5, comp["name"],
                ha='center', va='center', fontsize=10, fontweight='bold')

        ax.text(comp["x"] + comp["width"]/2, 0.5 - height/2 - 0.05, str(comp["shape"]),
                ha='center', va='top', fontsize=8)

        # Add connecting arrows
        if i > 0:
            prev = components[i-1]
            ax.annotate("",
                        xy=(comp["x"], 0.5),
                        xytext=(prev["x"] + prev["width"], 0.5),
                        arrowprops=dict(arrowstyle="->", color='black'))

    # Label axes
    ax.text(0.5, 0.95, "Convolutional Neural Network Architecture",
           ha='center', va='center', fontsize=14, fontweight='bold')

    # Set limits
    ax.set_xlim(0, 1)
    ax.set_ylim(0, 1)
    ax.axis('off')
```

```
plt.tight_layout()  
return fig
```

CNNs use specialized layers designed for processing grid-like data (e.g., images):

- **Convolutional Layers:** Apply filters to detect local patterns
 - Parameters: filter size, stride, padding, dilation
 - Properties: weight sharing, translation invariance
- **Pooling Layers:** Downsample feature maps
 - Types: max pooling, average pooling, global pooling
 - Purpose: Reduce dimensions, introduce invariance
- **Feature Hierarchy:** Early layers detect edges and textures; later layers detect complex shapes and objects

Notable CNN architectures include:

- **LeNet-5:** First successful CNN architecture
- **AlexNet:** Breakthrough in image classification (2012)
- **VGG:** Standardized architecture with small filters
- **Inception/GoogLeNet:** Parallel filter operations at different scales
- **ResNet:** Introduced residual connections, enabling extremely deep networks

10.4.2 Residual Networks

```
def visualize_residual_block():
    """Visualize a residual block from ResNet."""
    fig, ax = plt.subplots(figsize=(8, 6))

    # Draw main path
    # Input
    rect_input = plt.Rectangle((0.3, 0.8), 0.4, 0.1, facecolor='lightblue', edgecolor='black')
    ax.add_patch(rect_input)
    ax.text(0.5, 0.85, "Input", ha='center', va='center')

    # Conv 1
    rect_conv1 = plt.Rectangle((0.3, 0.65), 0.4, 0.1, facecolor='lightgreen', edgecolor='black')
    ax.add_patch(rect_conv1)
    ax.text(0.5, 0.7, "Conv 3×3", ha='center', va='center')

    # Batch Norm 1
    rect_bn1 = plt.Rectangle((0.3, 0.55), 0.4, 0.05, facecolor='lightyellow', edgecolor='black')
    ax.add_patch(rect_bn1)
    ax.text(0.5, 0.575, "BatchNorm", ha='center', va='center')

    # ReLU 1
    rect_relu1 = plt.Rectangle((0.3, 0.5), 0.4, 0.05, facecolor='lightpink', edgecolor='black')
    ax.add_patch(rect_relu1)
    ax.text(0.5, 0.525, "ReLU", ha='center', va='center')

    # Conv 2
    rect_conv2 = plt.Rectangle((0.3, 0.35), 0.4, 0.1, facecolor='lightgreen', edgecolor='black')
    ax.add_patch(rect_conv2)
    ax.text(0.5, 0.4, "Conv 3×3", ha='center', va='center')

    # Batch Norm 2
    rect_bn2 = plt.Rectangle((0.3, 0.25), 0.4, 0.05, facecolor='lightyellow', edgecolor='black')
    ax.add_patch(rect_bn2)
    ax.text(0.5, 0.275, "BatchNorm", ha='center', va='center')

    # Addition
    circle_add = plt.Circle((0.5, 0.15), 0.05, facecolor='white', edgecolor='black')
    ax.add_patch(circle_add)
    ax.text(0.5, 0.15, "+", ha='center', va='center', fontsize=15, fontweight='bold')

    # Output
    rect_output = plt.Rectangle((0.3, 0.05), 0.4, 0.05, facecolor='lightblue', edgecolor='black')
    ax.add_patch(rect_output)
    ax.text(0.5, 0.075, "Output", ha='center', va='center')

    # ReLU (final)
    rect_relu2 = plt.Rectangle((0.3, 0), 0.4, 0.05, facecolor='lightpink', edgecolor='black')
    ax.add_patch(rect_relu2)
    ax.text(0.5, 0.025, "ReLU", ha='center', va='center')

    # Draw shortcut path
```

```

ax.plot([0.5, 0.7, 0.7, 0.5], [0.8, 0.8, 0.15, 0.15], 'r-', linewidth=2)
ax.text(0.7, 0.45, "Shortcut Connection", ha='center', va='center', rotation=

# Draw arrows
arrow_props = dict(arrowstyle='->', color='black', linewidth=1.5)
ax.annotate('', xy=(0.5, 0.65), xytext=(0.5, 0.7), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.55), xytext=(0.5, 0.6), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.5), xytext=(0.5, 0.55), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.35), xytext=(0.5, 0.4), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.25), xytext=(0.5, 0.3), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.15), xytext=(0.5, 0.2), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0.05), xytext=(0.5, 0.1), arrowprops=arrow_props)
ax.annotate('', xy=(0.5, 0), xytext=(0.5, 0.05), arrowprops=arrow_props)

# Title
ax.text(0.5, 0.95, "ResNet Block", ha='center', va='center', fontsize=14, fon

# Set limits
ax.set_xlim(0, 1)
ax.set_ylim(0, 1)
ax.axis('off')

plt.tight_layout()
return fig

```

Residual Networks (ResNets) introduced skip connections to address the degradation problem in very deep networks:

$$\mathbf{y} = F(\mathbf{x}, \{W_i\}) + \mathbf{x}$$

where F represents the residual mapping and \mathbf{x} is the identity shortcut connection.

Benefits include:

- Easier optimization (shortcuts provide gradient highways)
- Better gradient flow in very deep networks
- Stabilized training
- State-of-the-art performance on many tasks

Variants of residual connections include:

- **Pre-activation ResNet:** Improved ordering of batch normalization and activation
- **ResNeXt:** Grouped convolutions for increased width
- **DenseNet:** Dense connections between all layers in a block
- **EfficientNet:** Balanced network depth, width, and resolution scaling

10.4.3 Normalization Techniques

```
def compare_normalizations():
    """Visualize differences between normalization techniques."""
    # Create toy feature maps (B, C, H, W)
    np.random.seed(42)
    batch_size = 4
    channels = 3
    height = 4
    width = 4

    # Create feature maps with different distributions per channel
    features = np.zeros((batch_size, channels, height, width))

    # Channel 0: Normal distribution with mean 10, std 5
    features[:, 0, :, :] = np.random.normal(10, 5, (batch_size, height, width))

    # Channel 1: Normal distribution with mean 0, std 1
    features[:, 1, :, :] = np.random.normal(0, 1, (batch_size, height, width))

    # Channel 2: Normal distribution with mean -5, std 3
    features[:, 2, :, :] = np.random.normal(-5, 3, (batch_size, height, width))

    # Apply different normalizations (simplified implementations)

    # Batch Normalization (normalize across batch, per channel)
    batch_norm = np.zeros_like(features)
    for c in range(channels):
        mean = np.mean(features[:, c, :, :])
        std = np.std(features[:, c, :, :])
        batch_norm[:, c, :, :] = (features[:, c, :, :] - mean) / (std + 1e-5)

    # Layer Normalization (normalize across channels, per sample)
    layer_norm = np.zeros_like(features)
    for b in range(batch_size):
        mean = np.mean(features[b, :, :, :])
        std = np.std(features[b, :, :, :])
        layer_norm[b, :, :, :] = (features[b, :, :, :] - mean) / (std + 1e-5)

    # Instance Normalization (normalize across spatial dims, per sample and channel)
    instance_norm = np.zeros_like(features)
    for b in range(batch_size):
        for c in range(channels):
            mean = np.mean(features[b, c, :, :])
            std = np.std(features[b, c, :, :])
            instance_norm[b, c, :, :] = (features[b, c, :, :] - mean) / (std + 1e-5)

    # Group Normalization (normalize across spatial dims and channel groups, per group)
    group_size = 1 # 1 group with 3 channels for this example
    group_norm = np.zeros_like(features)
    for b in range(batch_size):
        for g in range(0, channels, group_size):
            mean = np.mean(features[b, g:g+group_size, :, :])
```

```

        std = np.std(features[b, g:g+group_size, :, :])
        group_norm[b, g:g+group_size, :, :] = (features[b, g:g+group_size, :, :])

# Create visualizations
fig, axes = plt.subplots(2, 3, figsize=(15, 10))

# Helper function to visualize feature map distributions
def plot_distributions(ax, data, title):
    flat_data_by_channel = [data[:, c, :, :].flatten() for c in range(channel)]
    for c, channel_data in enumerate(flat_data_by_channel):
        ax.hist(channel_data, bins=20, alpha=0.7, label=f'Channel {c}')
    ax.set_title(title)
    ax.grid(True, alpha=0.3)
    ax.legend()

# Plot distributions
plot_distributions(axes[0, 0], features, 'Original Features')
plot_distributions(axes[0, 1], batch_norm, 'Batch Normalization')
plot_distributions(axes[0, 2], layer_norm, 'Layer Normalization')
plot_distributions(axes[1, 0], instance_norm, 'Instance Normalization')
plot_distributions(axes[1, 1], group_norm, 'Group Normalization')

# Diagram of normalization dimensions
axes[1, 2].axis('off')
axes[1, 2].text(0.5, 0.9, 'Normalization Dimensions', ha='center', fontsize=14)
axes[1, 2].text(0.5, 0.75, 'BatchNorm: Normalize across (N, H, W)', ha='center')
axes[1, 2].text(0.5, 0.65, 'LayerNorm: Normalize across (C, H, W)', ha='center')
axes[1, 2].text(0.5, 0.55, 'InstanceNorm: Normalize across (H, W)', ha='center')
axes[1, 2].text(0.5, 0.45, 'GroupNorm: Normalize across (G, H, W)', ha='center')
axes[1, 2].text(0.5, 0.3, 'N: Batch size, C: Channels', ha='center')
axes[1, 2].text(0.5, 0.2, 'H: Height, W: Width, G: Group', ha='center')

plt.tight_layout()
return fig

```

Different normalization techniques stabilize training by normalizing activations:

- **Batch Normalization:** Normalizes across the batch dimension
 - Pros: Very effective, improves training speed
 - Cons: Batch size dependent, less effective with small batches
 - $\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$
- **Layer Normalization:** Normalizes across all features for each sample
 - Pros: Batch size independent, good for recurrent networks
 - Cons: May not work well for CNNs
 - $\hat{x}_i = \frac{x_i - \mu_L}{\sqrt{\sigma_L^2 + \epsilon}}$

- **Instance Normalization:** Normalizes across spatial dimensions for each channel and sample
 - Pros: Effective for style transfer, independent of batch size
 - Cons: Loses statistical information about the dataset
 - $\hat{x}_{ijk} = \frac{x_{ijk} - \mu_{ij}}{\sqrt{\sigma_{ij}^2 + \epsilon}}$
- **Group Normalization:** Normalizes across groups of channels
 - Pros: Batch size independent, works well for smaller batches
 - Cons: Group size is a hyperparameter to tune
 - $\hat{x}_{ijg} = \frac{x_{ijg} - \mu_{ig}}{\sqrt{\sigma_{ig}^2 + \epsilon}}$

Choosing the right normalization technique depends on the architecture, task, and computational constraints.

10.4.4 Activation Functions

Modern activation functions improve on traditional ones like sigmoid and tanh:

- **ReLU:** Most common, but suffers from dying neurons
- **Leaky ReLU:** Prevents dying neurons with a small slope for negative inputs
- **Parametric ReLU (PReLU):** Learns the slope parameter during training
- **ELU (Exponential Linear Unit):** Smooth negative values with an exponential curve
- **GELU (Gaussian Error Linear Unit):** Used in transformers, approximates $x \cdot \Phi(x)$
- **Swish/SiLU:** Self-gated activation $x \cdot \sigma(x)$, often outperforms ReLU
- **Mish:** Smooth alternative to Swish with better performance

10.5 Biological Parallels in Deep Learning

Deep learning draws significant inspiration from neuroscience, though the connections are often overlooked in technical discussions. This section explores the parallels between neural networks and biological neural systems.

Biological Inspirations in Deep Learning

Parallels between neuroscience and artificial neural networks

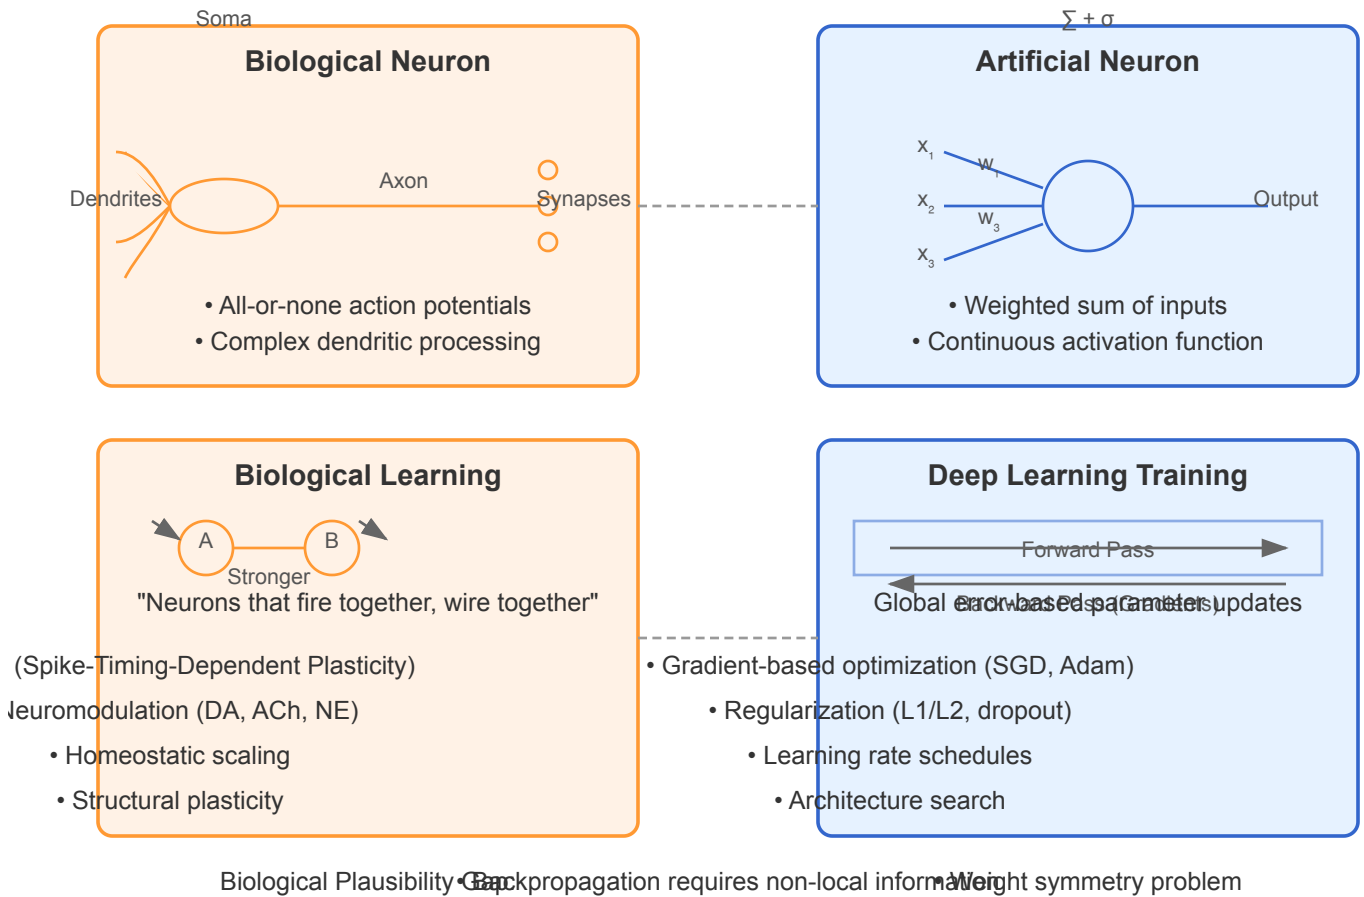


Figure 10.5: Comparison of biological and artificial neurons, highlighting similarities and differences in structure and learning mechanisms.

10.5.1 Neural Architectures and Brain Organization

While artificial neural networks are highly simplified compared to biological neurons, several architectural principles are shared:

- **Hierarchical Processing:** Both biological visual systems and CNNs process information in a hierarchical manner, with early layers detecting simple features and deeper layers representing more complex patterns.
- **Recurrent Connections:** Recurrent neural networks parallel the recurrent connectivity in cortical circuits, allowing for temporal processing and memory.
- **Attention Mechanisms:** Neural attention mechanisms are inspired by biological attention systems that selectively focus computational resources on relevant input features.

```

def biological_vs_artificial_neurons():
    """Compare biological and artificial neurons."""
    # Create a simple diagram
    fig, ax = plt.subplots(2, 1, figsize=(10, 8))

    # Biological neuron (simplified)
    ax[0].set_title("Biological Neuron")
    ax[0].axis('off')
    ax[0].text(0.1, 0.7, "Dendrites", fontsize=12)
    ax[0].text(0.5, 0.7, "Soma", fontsize=12)
    ax[0].text(0.8, 0.7, "Axon", fontsize=12)
    ax[0].text(0.9, 0.4, "Synapses", fontsize=12)

    # Draw simplified neuron
    ax[0].plot([0.1, 0.4], [0.5, 0.5], 'k-', linewidth=2) # Dendrite
    ax[0].plot([0.2, 0.35], [0.6, 0.5], 'k-', linewidth=2) # Dendrite
    ax[0].plot([0.3, 0.4], [0.4, 0.5], 'k-', linewidth=2) # Dendrite
    circle = plt.Circle((0.5, 0.5), 0.1, fill=True, color='lightgray')
    ax[0].add_patch(circle) # Soma
    ax[0].plot([0.6, 0.9], [0.5, 0.5], 'k-', linewidth=2) # Axon
    ax[0].plot([0.9, 0.95], [0.5, 0.3], 'k-', linewidth=1) # Synapse
    ax[0].plot([0.9, 0.95], [0.5, 0.5], 'k-', linewidth=1) # Synapse
    ax[0].plot([0.9, 0.95], [0.5, 0.7], 'k-', linewidth=1) # Synapse

    # Artificial neuron
    ax[1].set_title("Artificial Neuron")
    ax[1].axis('off')
    ax[1].text(0.05, 0.6, "x1", fontsize=12)
    ax[1].text(0.05, 0.5, "x2", fontsize=12)
    ax[1].text(0.05, 0.3, "xn", fontsize=12)
    ax[1].text(0.3, 0.7, "w1", fontsize=10)
    ax[1].text(0.3, 0.55, "w2", fontsize=10)
    ax[1].text(0.3, 0.3, "wn", fontsize=10)
    ax[1].text(0.4, 0.2, "Bias", fontsize=10)
    ax[1].text(0.5, 0.5, "Σ", fontsize=18)
    ax[1].text(0.7, 0.5, "σ", fontsize=16)
    ax[1].text(0.9, 0.5, "Output", fontsize=12)

    # Draw artificial neuron
    ax[1].plot([0.1, 0.4], [0.6, 0.5], 'k-', linewidth=1) # Input 1
    ax[1].plot([0.1, 0.4], [0.5, 0.5], 'k-', linewidth=1) # Input 2
    ax[1].plot([0.1, 0.4], [0.3, 0.5], 'k-', linewidth=1) # Input n
    ax[1].plot([0.4, 0.4], [0.3, 0.5], 'k:', linewidth=1) # ...
    ax[1].plot([0.4, 0.5], [0.3, 0.5], 'k-', linewidth=1) # Input n connection
    circle1 = plt.Circle((0.5, 0.5), 0.05, fill=True, color='lightgray')
    ax[1].add_patch(circle1) # Summation
    ax[1].plot([0.55, 0.65], [0.5, 0.5], 'k-', linewidth=1) # To activation
    circle2 = plt.Circle((0.7, 0.5), 0.05, fill=True, color='lightgray')
    ax[1].add_patch(circle2) # Activation
    ax[1].plot([0.75, 0.85], [0.5, 0.5], 'k-', linewidth=1) # Output

```

```
plt.tight_layout()  
return fig
```

10.5.2 Learning Mechanisms

The brain employs various learning mechanisms that have counterparts in deep learning:

- **Hebbian Learning vs. Backpropagation:** Hebbian learning ("neurons that fire together, wire together") is a local learning rule, while backpropagation propagates errors globally. Recent research explores biologically plausible alternatives to backpropagation, such as target propagation and feedback alignment.
- **Neuromodulation vs. Adaptive Learning Rates:** Neuromodulatory systems in the brain (dopamine, acetylcholine, etc.) regulate plasticity and learning, similar to how adaptive learning rate methods (Adam, RMSProp) modulate weight updates.
- **Homeostatic Plasticity vs. Regularization:** The brain employs homeostatic mechanisms to maintain stability, paralleling regularization techniques like weight decay and normalization in artificial networks.

10.5.3 Credit Assignment Problem

Both biological and artificial systems face the fundamental problem of credit assignment: determining which components contributed to an outcome.


```

def credit_assignment_comparison():
    """Compare credit assignment in biological and artificial systems."""
    fig, ax = plt.subplots(1, 2, figsize=(12, 5))

    # Biological credit assignment
    ax[0].set_title("Biological Credit Assignment")
    ax[0].axis('off')

    # Create a simple network diagram
    pos = {'A': (0.2, 0.8), 'B': (0.5, 0.8), 'C': (0.8, 0.8),
           'D': (0.2, 0.5), 'E': (0.5, 0.5), 'F': (0.8, 0.5),
           'G': (0.2, 0.2), 'H': (0.5, 0.2), 'I': (0.8, 0.2)}

    # Draw nodes
    for node, position in pos.items():
        circle = plt.Circle(position, 0.05, fill=True,
                             color='lightgray' if node not in ['A', 'I'] else 'lightb
        ax[0].add_patch(circle)
        ax[0].text(position[0], position[1], node,
                    ha='center', va='center', fontsize=12)

    # Draw edges
    edges = [('A', 'D'), ('A', 'E'), ('B', 'D'), ('B', 'E'), ('B', 'F'),
             ('C', 'E'), ('C', 'F'), ('D', 'G'), ('D', 'H'),
             ('E', 'G'), ('E', 'H'), ('E', 'I'), ('F', 'H'), ('F', 'I'),
             ('G', 'H'), ('G', 'I'), ('H', 'I')]

    for edge in edges:
        ax[0].plot([pos[edge[0]][0], pos[edge[1]][0]],
                    [pos[edge[0]][1], pos[edge[1]][1]],
                    'k-', alpha=0.6, linewidth=1)

    # Highlight local feedback paths
    ax[0].plot([pos['I'][0], pos['F'][0]], [pos['I'][1], pos['F'][1]],
                'r-', alpha=0.7, linewidth=2)
    ax[0].plot([pos['F'][0], pos['C'][0]], [pos['F'][1], pos['C'][1]],
                'r-', alpha=0.7, linewidth=2)

    # Add text explanation
    ax[0].text(0.5, 0.02, "Relies on local feedback signals and reward modulation",
               ha='center', fontsize=11)

    # Artificial credit assignment
    ax[1].set_title("Artificial Credit Assignment (Backpropagation)")
    ax[1].axis('off')

    # Use the same network layout
    for node, position in pos.items():
        circle = plt.Circle(position, 0.05, fill=True,
                             color='lightgray' if node not in ['A', 'I'] else 'lightb
        ax[1].add_patch(circle)
        ax[1].text(position[0], position[1], node,
                    ha='center', va='center', fontsize=12)

```

```

# Draw forward pass edges
for edge in edges:
    ax[1].plot([pos[edge[0]][0], pos[edge[1]][0]],
               [pos[edge[0]][1], pos[edge[1]][1]],
               'k-', alpha=0.6, linewidth=1)

# Draw backward pass (gradient flow)
backward_edges = [('I', 'F'), ('I', 'E'), ('I', 'H'), ('I', 'G'),
                  ('H', 'E'), ('H', 'D'), ('H', 'F'),
                  ('G', 'D'), ('G', 'E'),
                  ('F', 'C'), ('F', 'B'),
                  ('E', 'B'), ('E', 'A'), ('E', 'C'),
                  ('D', 'A'), ('D', 'B')]

for edge in backward_edges:
    ax[1].plot([pos[edge[0]][0], pos[edge[1]][0]],
               [pos[edge[0]][1], pos[edge[1]][1]],
               'r-', alpha=0.4, linewidth=1, linestyle='--')

# Add text explanation
ax[1].text(0.5, 0.02, "Propagates error backwards through entire network",
           ha='center', fontsize=11)

plt.tight_layout()
return fig

```

Backpropagation in artificial neural networks provides a mathematically precise solution to credit assignment but is not biologically plausible due to:

1. The need for symmetric weight matrices
2. Requiring precise storage of forward pass activations
3. Non-local weight updates

Neuroscience research explores alternatives like:

- **Three-factor Hebbian learning:** Combining pre/post-synaptic activity with a global modulation signal
- **Predictive coding:** Using prediction errors to drive learning
- **Feedback alignment:** Using random feedback weights for credit assignment

10.6 Modern Deep Learning Paradigms

Deep learning has evolved rapidly in recent years, with several important new paradigms emerging.

10.6.1 Self-Supervised Learning

Self-supervised learning has emerged as a powerful paradigm that leverages unlabeled data by creating “pseudo-labels” from the data itself.

```

def self_supervised_paradigms():
    """Illustrate different self-supervised learning approaches."""
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))

    # Masked Language Modeling
    axes[0, 0].set_title("Masked Language Modeling")
    axes[0, 0].axis('off')
    text = "The [MASK] jumped over the lazy dog."
    axes[0, 0].text(0.5, 0.7, text, ha='center', fontsize=12)
    axes[0, 0].text(0.5, 0.4, "↓", ha='center', fontsize=20)
    axes[0, 0].text(0.5, 0.2, "The fox jumped over the lazy dog.",
                    ha='center', fontsize=12, color='green')

    # Contrastive Learning
    axes[0, 1].set_title("Contrastive Learning")
    axes[0, 1].axis('off')
    # Draw an anchor image
    rect1 = plt.Rectangle((0.3, 0.6), 0.4, 0.3, fill=True, color='lightblue')
    axes[0, 1].add_patch(rect1)
    axes[0, 1].text(0.5, 0.75, "Anchor", ha='center')

    # Draw positive and negative examples
    rect2 = plt.Rectangle((0.1, 0.2), 0.3, 0.2, fill=True, color='lightblue')
    axes[0, 1].add_patch(rect2)
    axes[0, 1].text(0.25, 0.3, "Positive", ha='center')

    rect3 = plt.Rectangle((0.6, 0.2), 0.3, 0.2, fill=True, color='lightcoral')
    axes[0, 1].add_patch(rect3)
    axes[0, 1].text(0.75, 0.3, "Negative", ha='center')

    # Draw attraction/repulsion arrows
    axes[0, 1].arrow(0.4, 0.6, -0.1, -0.25, head_width=0.02, head_length=0.02,
                    fc='green', ec='green')
    axes[0, 1].arrow(0.6, 0.6, 0.1, -0.25, head_width=0.02, head_length=0.02,
                    fc='red', ec='red')

    # Autoregressive Prediction
    axes[1, 0].set_title("Autoregressive Prediction")
    axes[1, 0].axis('off')
    text = "The fox jumped over the"
    axes[1, 0].text(0.5, 0.7, text, ha='center', fontsize=12)
    axes[1, 0].text(0.5, 0.4, "↓", ha='center', fontsize=20)
    axes[1, 0].text(0.5, 0.2, "lazy dog.",
                    ha='center', fontsize=12, color='green')

    # Rotation/Colorization
    axes[1, 1].set_title("Image Restoration")
    axes[1, 1].axis('off')

    # Draw a grayscale or corrupted image
    rect4 = plt.Rectangle((0.1, 0.6), 0.3, 0.3, fill=True, color='lightgray')
    axes[1, 1].add_patch(rect4)
    axes[1, 1].text(0.25, 0.5, "Corrupted Input", ha='center')

```

```
# Draw arrow
axes[1, 1].text(0.5, 0.7, ">", ha='center', fontsize=20)

# Draw restored image
rect5 = plt.Rectangle((0.6, 0.6), 0.3, 0.3, fill=True, color='lightblue')
axes[1, 1].add_patch(rect5)
axes[1, 1].text(0.75, 0.5, "Restored Image", ha='center')

plt.tight_layout()
return fig
```

Key self-supervised paradigms include:

- **Masked Language/Image Modeling:** Predicting masked tokens from surrounding context
- **Contrastive Learning:** Learning to distinguish between similar and dissimilar examples
- **Autoregressive Prediction:** Predicting next elements in a sequence
- **Data Restoration:** Reconstructing corrupted or modified versions of the input

Self-supervised learning has enabled state-of-the-art results across domains with limited labeled data, and forms the foundation of modern foundation models.

10.6.2 Foundation Models and Scaling Laws

Recent years have witnessed the emergence of foundation models: large-scale models pre-trained on vast amounts of data that can be adapted to various downstream tasks.

```

def scaling_laws():
    """Visualize scaling laws in deep learning."""
    fig, ax = plt.subplots(figsize=(10, 6))

    # Log scales
    ax.set_xscale('log')
    ax.set_yscale('log')

    # Compute data points
    x = np.logspace(0, 4, 100)
    y_params = 0.5 * x**(-0.3) # Performance improves with model size
    y_data = 0.5 * x**(-0.25) # Performance improves with data
    y_compute = 0.5 * x**(-0.2) # Performance improves with compute

    # Plot scaling curves
    ax.plot(x, y_params, 'b-', label='Model Size Scaling')
    ax.plot(x, y_data, 'r-', label='Dataset Size Scaling')
    ax.plot(x, y_compute, 'g-', label='Compute Scaling')

    # Add scaling regimes markers
    ax.axvline(x=10, color='gray', linestyle='--', alpha=0.5)
    ax.axvline(x=1000, color='gray', linestyle='--', alpha=0.5)

    ax.text(5, 0.01, "Small\nModels", ha='right')
    ax.text(500, 0.01, "Medium\nModels", ha='center')
    ax.text(5000, 0.01, "Large\nModels", ha='left')

    # Labels
    ax.set_xlabel('Scale Factor (log)')
    ax.set_ylabel('Loss (log)')
    ax.set_title('Scaling Laws in Deep Learning')
    ax.legend()

    # Emergent abilities annotation
    ax.annotate('Emergent\nAbilities', xy=(1000, 0.05), xytext=(500, 0.15),
               arrowprops=dict(arrowstyle='->'))

    plt.grid(True, which="both", ls="--", alpha=0.2)
    return fig

```

Key insights from scaling research include:

- **Predictable Scaling Laws:** Model performance improves following power laws with increases in model size, data, and compute.
- **Emergent Abilities:** Beyond certain scale thresholds, models demonstrate qualitatively new capabilities not present in smaller models.
- **Transfer Learning Efficiency:** Large pre-trained models can be efficiently fine-tuned for downstream tasks with relatively little task-specific data.

Foundation models have transformed deep learning research and applications, with models like:

- **Large Language Models:** GPT, LLaMA, Claude
- **Vision-Language Models:** CLIP, DALL-E, Stable Diffusion
- **Multimodal Models:** GPT-4, Gemini

10.6.3 Loss Landscapes

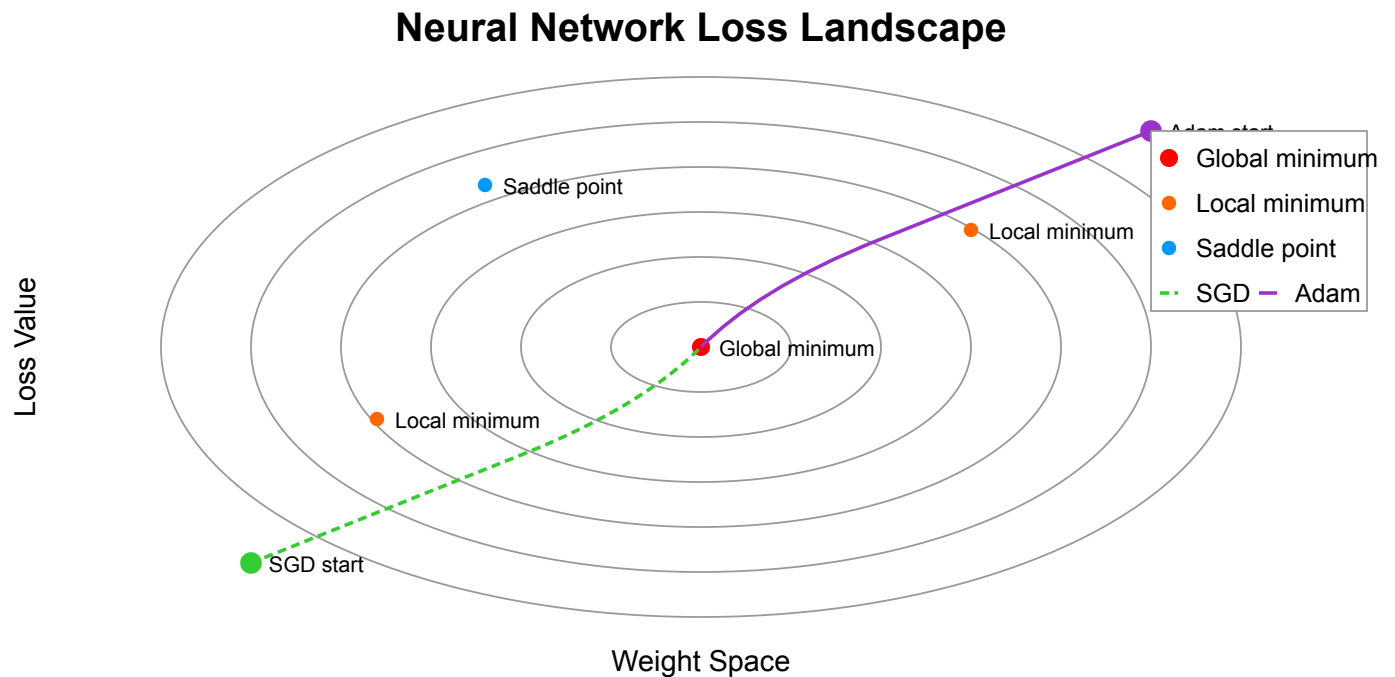


Figure 10.6: Visualization of neural network loss landscape showing the complex optimization surface with local minima, saddle points, and flat regions.

Loss landscapes in deep networks are complex, high-dimensional surfaces with many local minima, saddle points, and flat regions:

- **Local Minima:** Points where the loss is lower than all nearby points
- **Global Minimum:** The lowest possible loss value
- **Saddle Points:** Points with zero gradient but not minima (common in high dimensions)
- **Flat Regions:** Areas with very small gradients that slow training
- **Sharp Minima:** Minima with high curvature, often associated with poor generalization
- **Wide Minima:** Minima with low curvature, often associated with good generalization

Recent research suggests that most critical points in deep networks are saddle points rather than local minima, and that finding wide minima leads to better generalization.

10.5.2 Generalization Theory

Generalization is the ability of a model to perform well on unseen data:

- **Empirical Risk Minimization:** Minimizing loss on training data
- **Structural Risk Minimization:** Balancing empirical risk and model complexity
- **Regularization:** Constraining model complexity to improve generalization
- **VC Dimension:** Theoretical measure of model capacity
- **Rademacher Complexity:** Measure of a model's ability to fit random noise

Modern deep learning often violates classical generalization bounds because models can memorize random data yet still generalize well on real data. This paradox has led to new theories:

- **Flat Minima Hypothesis:** Models that find flat regions of the loss landscape generalize better
- **Implicit Regularization:** Optimization methods like SGD inherently bias toward simpler solutions
- **Neural Tangent Kernel:** Connects neural network training to kernel methods in the infinite-width limit

10.5.3 Double Descent Phenomenon

```
def double_descent_curve():
    """Visualize the double descent phenomenon."""
    # Model complexity (e.g., number of parameters)
    complexity = np.linspace(1, 100, 1000)

    # Critical complexity where model can perfectly fit training data
    critical_complexity = 40

    # Classical U-shaped risk curve
    classical_risk = 1.0 / (complexity + 0.1) + 0.02 * complexity

    # Double descent risk curve
    interpolation_peak = 5.0 * np.exp(-0.2 * (complexity - critical_complexity)**2)
    modern_risk = 1.0 / (complexity + 0.1) + 0.01 * np.exp(-0.05 * complexity) +

    # Training error (decreases monotonically)
    train_error = 2.0 / (1 + np.exp(0.1 * (complexity - critical_complexity))) -

    # Plot
    plt.figure(figsize=(10, 6))

    plt.plot(complexity, classical_risk, 'r--', label='Classical Theory (U-shape)')
    plt.plot(complexity, modern_risk, 'b-', label='Modern Observation (Double Des)')
    plt.plot(complexity, train_error, 'g-.', label='Training Error')

    # Mark interpolation threshold
    plt.axvline(x=critical_complexity, color='gray', linestyle=':', alpha=0.7)
    plt.text(critical_complexity + 1, 2.5, 'Interpolation Threshold', rotation=90)

    # Annotate regions
    plt.annotate('Underfitting', xy=(10, 1.2), xytext=(10, 2.0),
                arrowprops=dict(arrowstyle='->'))

    plt.annotate('Interpolation\nRegime', xy=(critical_complexity, 2.5), xytext=(
        critical_complexity + 10, 2.5),
                arrowprops=dict(arrowstyle='->'))

    plt.annotate('Modern Generalization', xy=(80, 0.5), xytext=(70, 1.5),
                arrowprops=dict(arrowstyle='->'))

    plt.xlabel('Model Complexity')
    plt.ylabel('Risk (Test Error)')
    plt.title('Double Descent Phenomenon')
    plt.legend()
    plt.grid(True, alpha=0.3)

    plt.tight_layout()
    return plt
```

The double descent phenomenon challenges the classical bias-variance tradeoff:

1. **Classical U-curve:** As model complexity increases, test error first decreases (reducing bias), then increases (increasing variance)
2. **Double Descent:** After the interpolation threshold (where training error reaches zero), test error can *decrease again* with increasing model complexity

This phenomenon helps explain why overparameterized deep networks (with more parameters than training examples) can still generalize well.

10.5.4 Neural Tangent Kernel

The Neural Tangent Kernel (NTK) is a theoretical tool for understanding neural network training:

- Connects neural networks to kernel methods
- Shows that in the infinite-width limit, neural networks behave like linear models in a fixed feature space
- Explains why wide networks train stably and generalize well
- Predicts training dynamics of wide networks

While primarily theoretical, NTK insights inform network initialization and architecture design.

10.6 Code Lab: Implementing a Neural Network from Scratch

Let's implement a simple neural network without using deep learning frameworks to understand the core concepts better:

```

def neural_network_from_scratch():
    """Implement a simple neural network from scratch."""
    # Define network architecture
    input_size = 2
    hidden_size = 3
    output_size = 1

    # Generate synthetic data
    np.random.seed(42)
    X = np.random.randn(100, input_size)
    # True function: XOR-like (non-linear)
    y = np.array([(x[0] > 0) != (x[1] > 0) for x in X]).reshape(-1, 1).astype(float)

    # Initialize weights and biases
    def init_params():
        np.random.seed(42)
        W1 = np.random.randn(input_size, hidden_size) * 0.1
        b1 = np.zeros((1, hidden_size))
        W2 = np.random.randn(hidden_size, output_size) * 0.1
        b2 = np.zeros((1, output_size))
        return {'W1': W1, 'b1': b1, 'W2': W2, 'b2': b2}

    # Activation functions
    def sigmoid(x):
        return 1 / (1 + np.exp(-x))

    def sigmoid_derivative(x):
        s = sigmoid(x)
        return s * (1 - s)

    # Forward pass
    def forward(X, params):
        W1, b1, W2, b2 = params['W1'], params['b1'], params['W2'], params['b2']

        # Hidden layer
        Z1 = X @ W1 + b1
        A1 = sigmoid(Z1)

        # Output layer
        Z2 = A1 @ W2 + b2
        A2 = sigmoid(Z2)

        cache = {'Z1': Z1, 'A1': A1, 'Z2': Z2, 'A2': A2, 'X': X}
        return A2, cache

    # Compute loss
    def compute_loss(A2, y):
        m = y.shape[0]
        loss = -np.sum(y * np.log(A2 + 1e-8) + (1 - y) * np.log(1 - A2 + 1e-8)) / m
        return loss

    # Backward pass
    def backward(cache, y, params):

```

```

m = y.shape[0]
W1, W2 = params['W1'], params['W2']
A1, A2 = cache['A1'], cache['A2']
X = cache['X']

# Output layer
dZ2 = A2 - y
dW2 = A1.T @ dZ2 / m
db2 = np.sum(dZ2, axis=0, keepdims=True) / m

# Hidden layer
dA1 = dZ2 @ W2.T
dZ1 = dA1 * sigmoid_derivative(cache['Z1'])
dW1 = X.T @ dZ1 / m
db1 = np.sum(dZ1, axis=0, keepdims=True) / m

gradients = {'dW1': dW1, 'db1': db1, 'dW2': dW2, 'db2': db2}
return gradients

# Update parameters
def update_params(params, gradients, learning_rate):
    params['W1'] -= learning_rate * gradients['dW1']
    params['b1'] -= learning_rate * gradients['db1']
    params['W2'] -= learning_rate * gradients['dW2']
    params['b2'] -= learning_rate * gradients['db2']
    return params

# Training loop
def train(X, y, hidden_size, learning_rate, epochs):
    # Initialize parameters
    params = init_params()

    # Track loss
    losses = []

    # Training iterations
    for i in range(epochs):
        # Forward pass
        A2, cache = forward(X, params)

        # Compute loss
        loss = compute_loss(A2, y)
        losses.append(loss)

        # Backward pass
        gradients = backward(cache, y, params)

        # Update parameters
        params = update_params(params, gradients, learning_rate)

        # Print loss every 1000 epochs
        if i % 1000 == 0:
            print(f"Epoch {i}, Loss: {loss:.4f}")

```

```

        return params, losses

# Train the network
params, losses = train(X, y, hidden_size, learning_rate=0.5, epochs=5000)

# Visualize the results
def visualize_results():
    # Visualization of training progress
    plt.figure(figsize=(12, 5))

    # Loss curve
    plt.subplot(1, 2, 1)
    plt.plot(losses)
    plt.title('Training Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.grid(True)

    # Decision boundary
    plt.subplot(1, 2, 2)

    # Create a mesh grid
    h = 0.01
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))

    # Make predictions on the mesh grid
    Z, _ = forward(np.c_[xx.ravel(), yy.ravel()], params)
    Z = Z.reshape(xx.shape)

    # Plot decision boundary
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral, alpha=0.8)

    # Plot training examples
    plt.scatter(X[:, 0], X[:, 1], c=y.ravel(), cmap=plt.cm.Spectral, edgecolor='k')
    plt.title('Decision Boundary')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')

    plt.tight_layout()
    return plt

return visualize_results()

```

This implementation demonstrates the core components of neural networks:

1. **Forward Propagation:** Computing activations through the network
2. **Loss Calculation:** Measuring how far predictions are from targets
3. **Backward Propagation:** Computing gradients for each parameter

4. **Parameter Updates:** Adjusting weights and biases using gradients

While modern deep learning frameworks automate these steps, understanding the underlying mechanics is crucial for debugging, customization, and optimization.

10.7 Take-aways

- **Framework Abstraction:** Modern deep learning frameworks like PyTorch and TensorFlow abstract low-level details, allowing researchers to focus on architecture design and experimentation.
- **Training Stability:** Achieving stable training requires careful attention to initialization, normalization, learning rates, and gradient flow.
- **Regularization Importance:** Regularization techniques critically impact model generalization, with techniques like dropout, batch normalization, and weight decay combining for best results.
- **Architectural Innovations:** Advances like residual connections enable training of extremely deep networks by mitigating gradient flow issues.
- **Optimization Challenges:** Deep learning optimization remains challenging due to non-convex loss landscapes, saddle points, and the need to escape poor local minima.
- **Empirical Focus:** Despite theoretical progress, deep learning remains heavily empirical, with practical techniques often preceding theoretical understanding.

! Knowledge Connections

Looking Back

- **Chapter 1 (Introduction):** The backpropagation algorithm introduced in section 1.1.3 is explored in depth here, showing how it enables training of complex neural networks.
- **Chapter 2 (Neuroscience Foundations):** The biological learning mechanisms described in section 2.3 provide an interesting contrast to the optimization algorithms covered in this chapter.
- **Chapter 7 (Information Theory):** The information bottleneck principle (section 7.6) provides theoretical insights into how deep networks compress information through layers.
- **Chapter 9 (ML Foundations):** The basic learning algorithms from Chapter 9 are extended here to handle deep architectures and large-scale training.

Looking Forward

- **Chapter 11 (Sequence Models):** The optimization techniques learned here will be applied to the specialized architectures for sequential data.
- **Chapter 12 (Large Language Models):** The scaling laws and foundation model concepts introduced in section 10.6.2 become critical for understanding LLM training and capabilities.
- **Chapter 14 (Future Directions):** The biological parallels in deep learning (section 10.5) point to neuromorphic approaches that may shape future AI systems.

10.8 Further Reading & Media

- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
<http://www.deeplearningbook.org>
- Smith, L. N. (2018). "A disciplined approach to neural network hyper-parameters: Part 1 - learning rate, batch size, momentum, and weight decay." *arXiv preprint arXiv:1803.09820*.
- Li, H., Xu, Z., Taylor, G., Studer, C., & Goldstein, T. (2018). "Visualizing the Loss Landscape of Neural Nets." *Advances in Neural Information Processing Systems*.
- Loshchilov, I., & Hutter, F. (2019). "Decoupled Weight Decay Regularization." *International Conference on Learning Representations*.

- Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2021). "Understanding deep learning (still) requires rethinking generalization." *Communications of the ACM*, 64(3), 107-115.
- He, K., Zhang, X., Ren, S., & Sun, J. (2016). "Deep Residual Learning for Image Recognition." *IEEE Conference on Computer Vision and Pattern Recognition*.
- Jacot, A., Gabriel, F., & Hongler, C. (2018). "Neural Tangent Kernel: Convergence and Generalization in Neural Networks." *Advances in Neural Information Processing Systems*.
- Ioffe, S., & Szegedy, C. (2015). "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift." *International Conference on Machine Learning*.
- 3Blue1Brown YouTube Series: "Neural Networks" https://www.youtube.com/playlist?list=PLZHQObOWTQDNU6R1_67000Dx_ZCJB-3pi